

# *Unsupervised Machine Learning for Customer Market Segmentation*

Guided Project completed by Suhaimi William Chan

Instructor: Ryan Ahmed

## Course Objectives

In this course, we are going to focus on the following learning objectives:

1. *Understand how to leverage the power of machine learning to transform marketing departments and perform customer segmentation*
2. *Apply Python libraries to import and visualize dataset images.*
3. *Understand the theory and intuition behind k-means clustering machine learning algorithm*
4. *Learn how to obtain the optimal number of clusters using the elbow method*
5. *Apply Scikit-Learn library to find the optimal number of clusters using elbow method*
6. *Apply k-means in Scikit-Learn to perform customer segmentation*
7. *Understand the theory and intuition behind Principal Component Analysis (PCA) algorithm*
8. *Apply Principal Component Analysis (PCA) technique to perform dimensionality reduction and data visualization*
9. *Compile and fit unsupervised machine learning models such as PCA and K-Means to training data*

# Project Structure

The hands on project on **Unsupervised Machine Learning for Customer Segmentation** is divided into following tasks:

Task 1: Understand the problem statement and business case

Task 2: Import libraries and datasets

Task 3: Visualize and explore datasets

Task 4: Understand the theory and intuition behind k-means clustering machine learning algorithm

Task 5: Learn how to obtain the optimal number of clusters using the elbow method

Task 6: Use Scikit-Learn library to find the optimal number of clusters using elbow method

Task 7: Apply k-means using Scikit-Learn to perform customer segmentation

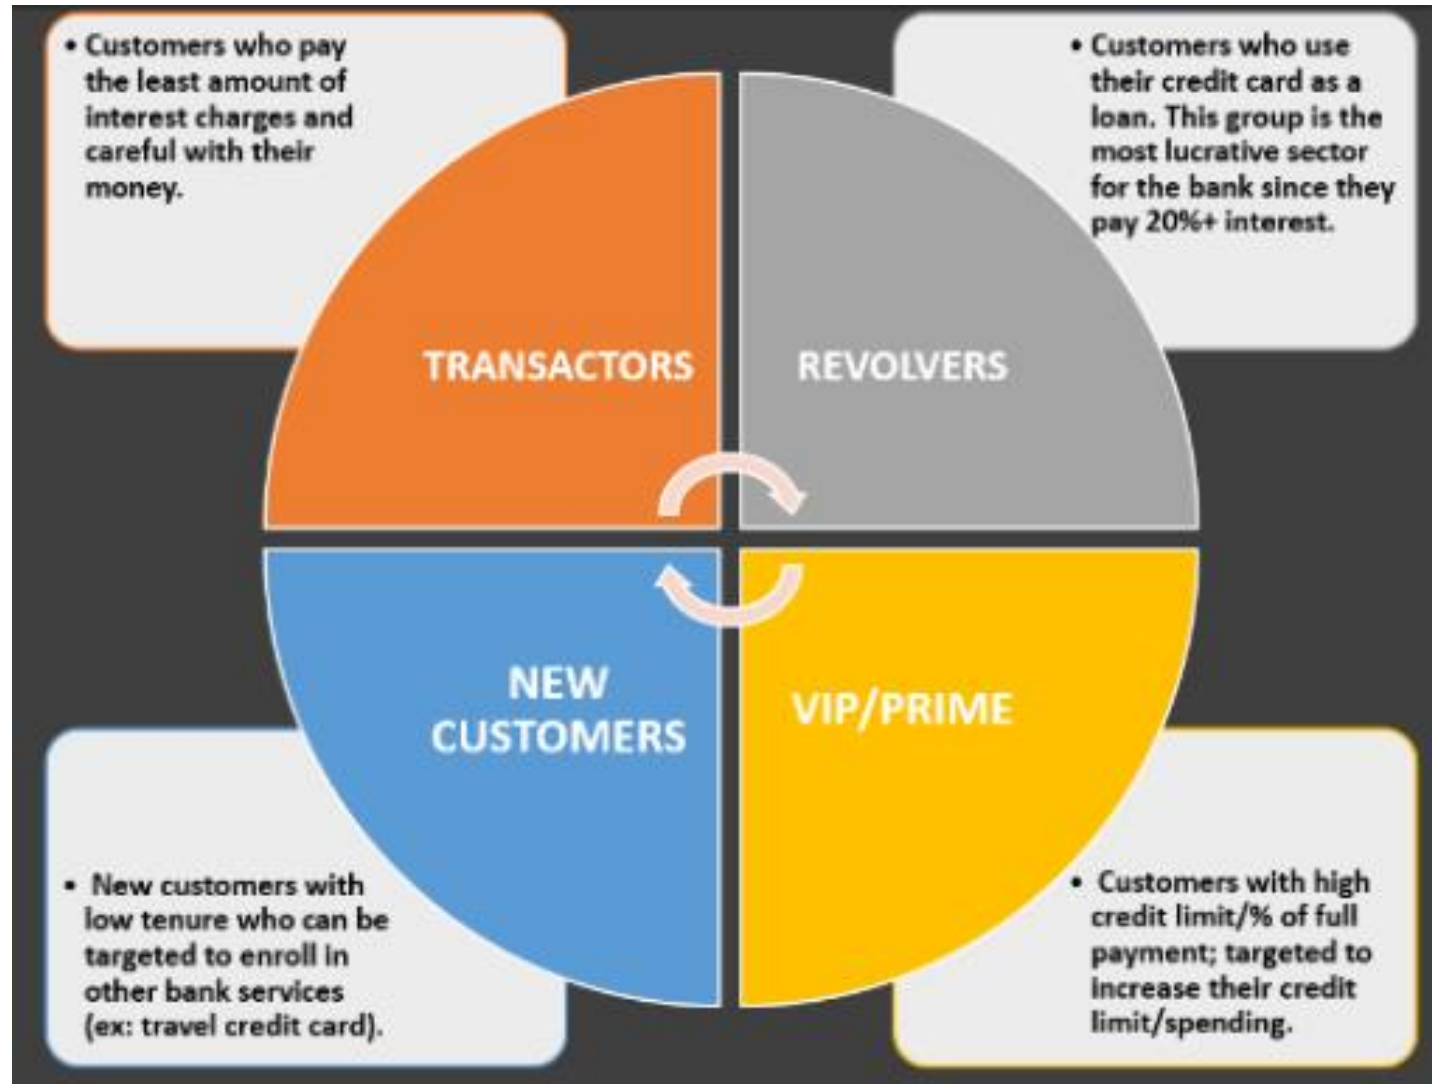
Task 8: Apply Principal Component Analysis (PCA) technique to perform dimensionality reduction and data visualization

## TASK #1: UNDERSTAND THE PROBLEM STATEMENT AND BUSINESS CASE

- In this project, you have been hired as a data scientist at a bank and you have been provided with extensive data on the bank's customers for the past 6 months.
- Data includes transactions frequency, amount, tenure..etc.
- The bank marketing team would like to leverage AI/ML to launch a targeted marketing ad campaign that is tailored to specific group of customers.
- In order for this campaign to be successful, the bank has to divide its customers into at least 3 distinctive groups.
- This process is known as "marketing segmentation" and it crucial for maximizing marketing campaign conversion rate.



- Data Source: <https://www.kaggle.com/ajunbhasin2013/ccdata>
- Photo Credit: <https://www.pexels.com/photo/101172/marketing-customer-polaroid-center-presentation-online-board-target-economy>



## TASK #2: IMPORT LIBRARIES AND DATASETS

```
In [ ]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler, normalize
from sklearn.cluster import KMeans
from sklearn.decomposition import PCA
from jupyterthemes import jtplot
jtplot.style(theme='monokai', context='notebook', ticks=True, grid=False)
# setting the style of the notebook to be monokai theme
# this line of code is important to ensure that we are able to see the x and y axes clearly
# If you don't run this code line, you will notice that the xlabel and ylabel on any plot is black on black
```

```
In [ ]: # You have to include the full link to the csv file containing your dataset
creditcard_df = pd.read_csv('Marketing_data.csv')

# CUSTID: Identification of Credit Card holder
# BALANCE: Balance amount left in customer's account to make purchases
# BALANCE_FREQUENCY: How frequently the Balance is updated, score between 0 and 1 (1 = frequently updated, 0 = not frequently updated)
# PURCHASES: Amount of purchases made from account
# ONEOFFPURCHASES: Maximum purchase amount done in one-go
```

```
# CUSTID: Identification of Credit Card holder
# BALANCE: Balance amount left in customer's account to make purchases
# BALANCE_FREQUENCY: How frequently the Balance is updated, score between 0 and 1 (1 = frequently updated, 0 = not frequently updated)
# PURCHASES: Amount of purchases made from account
# ONEOFFPURCHASES: Maximum purchase amount done in one-go
# INSTALLMENTS_PURCHASES: Amount of purchase done in installment
# CASH_ADVANCE: Cash in advance given by the user
# PURCHASES_FREQUENCY: How frequently the Purchases are being made, score between 0 and 1 (1 = frequently purchased, 0 = not frequently purchased)
# ONEOFF_PURCHASES_FREQUENCY: How frequently Purchases are happening in one-go (1 = frequently purchased, 0 = not frequently purchased)
# PURCHASES_INSTALLMENTS_FREQUENCY: How frequently purchases in installments are being done (1 = frequently, 0 = not frequently)
# CASH_ADVANCE_FREQUENCY: How frequently the cash in advance being paid
# CASH_ADVANCE_TRX: Number of Transactions made with "Cash in Advance"
# PURCHASES_TRX: Number of purchase transactions made
# CREDIT_LIMIT: Limit of Credit Card for user
# PAYMENTS: Amount of Payment done by user
# MINIMUM_PAYMENTS: Minimum amount of payments made by user
# PRC_FULL_PAYMENT: Percent of full payment paid by user
# TENURE: Tenure of credit card service for user
```



```
In [3]: creditcard_df
```

	CUST_ID	BALANCE	BALANCE_FREQUENCY	PURCHASES	ONEOFF_PURCHASES	INSTALLMENTS_PURCHASES	CASH_ADVANCE
0	C10001	40.900749	0.818182	95.40	0.00	95.40	0.000000
1	C10002	3202.467416	0.909091	0.00	0.00	0.00	6442.000000
2	C10003	2495.148862	1.000000	773.17	773.17	0.00	0.000000
3	C10004	1666.670542	0.636364	1499.00	1499.00	0.00	205.170000
4	C10005	817.714335	1.000000	16.00	16.00	0.00	0.000000
...	...	...	...	...	...	...	...
8945	C19186	28.493517	1.000000	291.12	0.00	291.12	0.000000
8946	C19187	19.183215	1.000000	300.00	0.00	300.00	0.000000

8950 rows x 18 columns

In [4]:

```
creditcard_df.info()
```

```
# Let's apply info and get additional insights on our dataframe
```

```
# 18 features with 8950 points
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 8950 entries, 0 to 8949
```

```
Data columns (total 18 columns):
```

#	Column	Non-Null Count	Dtype
0	CUST_ID	8950 non-null	object
1	BALANCE	8950 non-null	float64
2	BALANCE_FREQUENCY	8950 non-null	float64
3	PURCHASES	8950 non-null	float64
4	ONEOFF_PURCHASES	8950 non-null	float64
5	INSTALLMENTS_PURCHASES	8950 non-null	float64
6	CASH_ADVANCE	8950 non-null	float64
7	PURCHASES_FREQUENCY	8950 non-null	float64
8	ONEOFF_PURCHASES_FREQUENCY	8950 non-null	float64
9	PURCHASES_INSTALLMENTS_FREQUENCY	8950 non-null	float64
10	CASH_ADVANCE_FREQUENCY	8950 non-null	float64
11	CASH_ADVANCE_TRX	8950 non-null	int64
12	PURCHASES_TRX	8950 non-null	int64
13	CREDIT_LIMIT	8949 non-null	float64
14	PAYMENTS	8950 non-null	float64
15	MINIMUM_PAYMENTS	8637 non-null	float64
16	PRC_FULL_PAYMENT	8950 non-null	float64
17	TENURE	8950 non-null	int64

```
dtypes: float64(14), int64(3), object(1)
```

```
memory usage: 1.2+ MB
```



### MINI CHALLENGE #1:

- What is the average, minimum and maximum "BALANCE" amount?

```
In [8]: print('Average, min, max =', creditcard_df['BALANCE'].mean(),  
          creditcard_df['BALANCE'].min(), creditcard_df['BALANCE'].max())
```

```
Average, min, max = 1564.4748276781038 0.0 19043.13856
```

```
In [ ]: # Let's apply describe() and get more statistical insights on our dataframe  
        # Mean balance is $1564  
        # Balance frequency is frequently updated on average ~0.9  
        # Purchases average is $1000  
        # one off purchase average is ~$600  
        # Average purchases frequency is around 0.5  
        # average ONEOFF_PURCHASES_FREQUENCY, PURCHASES_INSTALLMENTS_FREQUENCY, and CASH_ADVANCE_FREQUENCY are gene  
        # Average credit limit ~ 4500  
        # Percent of full payment is 15%  
        # Average tenure is 11 years
```

```
In [7]: creditcard_df.describe()

# Let's apply describe() and get more statistical insights on our dataframe
# Mean balance is $1564
# Balance frequency is frequently updated on average ~0.9
# Purchases average is $1000
# one off purchase average is ~$600
# Average purchases frequency is around 0.5
# average ONEOFF_PURCHASES_FREQUENCY, PURCHASES_INSTALLMENTS_FREQUENCY, and CASH_ADVANCE_FREQUENCY are general
# Average credit limit ~ 4500
# Percent of full payment is 15%
# Average tenure is 11 years
```

	BALANCE	BALANCE_FREQUENCY	PURCHASES	ONEOFF_PURCHASES	INSTALLMENTS_PURCHASES	CASH_ADVANCE
count	8950.000000	8950.000000	8950.000000	8950.000000	8950.000000	8950.000000
mean	1564.474828	0.877271	1003.204834	592.437371	411.067645	978.871112
std	2081.531879	0.236904	2136.634782	1659.887917	904.338115	2097.163877
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	128.281915	0.888889	39.635000	0.000000	0.000000	0.000000
50%	873.385231	1.000000	361.280000	38.000000	89.000000	0.000000
75%	2054.140036	1.000000	1110.130000	577.405000	468.637500	1113.821139
max	19043.138560	1.000000	49039.570000	40761.250000	22500.000000	47137.211760

### MINI CHALLENGE #2:

- Obtain the features (row) of the customer who made the maximum "ONEOFF\_PURCHASES"
- Obtain the features of the customer who made the maximum cash advance transaction? how many cash advance transactions did that customer make? how often did he/she pay their bill?

```
In [8]: creditcard_df[creditcard_df['ONEOFF_PURCHASES'] == 40761.25]
```

	CUST_ID	BALANCE	BALANCE_FREQUENCY	PURCHASES	ONEOFF_PURCHASES	INSTALLMENTS_PURCHASES	CASH_ADVANCE
550	C10574	11547.52001	1.0	49039.57	40761.25	8278.32	558.16

```
In [10]: creditcard_df['CASH_ADVANCE'].max()
```

47137.211760000006

```
In [12]: creditcard_df[creditcard_df['CASH_ADVANCE'] == 47137.211760000006]
```

	INSTALLMENTS_FREQUENCY	CASH_ADVANCE_FREQUENCY	CASH_ADVANCE_TRX	PURCHASES_TRX	CREDIT_LIMIT	PAYMENTS
	1.0	123	21	19600.0	39048.59762	

```
In [12]: creditcard_df[creditcard_df['CASH_ADVANCE'] == 47137.211760000006]
```

	CASH_ADVANCE_TRX	PURCHASES_TRX	CREDIT_LIMIT	PAYMENTS	MINIMUM_PAYMENTS	PRC_FULL_PAYMENT	TENURE
123	21	19600.0	39048.59762	5394.173671	0.0	12	

## # TASK #3: VISUALIZE AND EXPLORE DATASET

In [ ]:

```
In [13]: # Let's see if we have any missing data, Luckily we don't have many!  
sns.heatmap(creditcard_df.isnull(), yticklabels = False, cbar = False, cmap="Blues")
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x2628cdcb508>

CUST_ID	
BALANCE	
BALANCE_FREQUENCY	
PURCHASES	
ONEOFF_PURCHASES	
INSTALLMENTS_PURCHASES	
CASH_ADVANCE	
PURCHASES_FREQUENCY	
ONEOFF_PURCHASES_FREQUENCY	
PURCHASES_INSTALLMENTS_FREQUENCY	
CASH_ADVANCE_FREQUENCY	
CASH_ADVANCE_TRX	
PURCHASES_TRX	
CREDIT_LIMIT	
PAYMENTS	
MINIMUM_PAYMENTS	
PRC_FULL_PAYMENT	
TENURE	

```
In [14]: creditcard_df.isnull().sum()
```

CUST_ID	0
BALANCE	0
BALANCE_FREQUENCY	0
PURCHASES	0
ONEOFF_PURCHASES	0
INSTALLMENTS_PURCHASES	0
CASH_ADVANCE	0
PURCHASES_FREQUENCY	0
ONEOFF_PURCHASES_FREQUENCY	0
PURCHASES_INSTALLMENTS_FREQUENCY	0
CASH_ADVANCE_FREQUENCY	0
CASH_ADVANCE_TRX	0
PURCHASES_TRX	0
CREDIT_LIMIT	1
PAYMENTS	0
MINIMUM_PAYMENTS	313
PRC_FULL_PAYMENT	0
TENURE	0
dtype:	int64



```
In [16]: # Fill up the missing elements with mean of the 'MINIMUM_PAYMENT'
creditcard_df.loc[(creditcard_df['MINIMUM_PAYMENTS'].isnull() == True),
                  'MINIMUM_PAYMENTS'] = creditcard_df['MINIMUM_PAYMENTS'].mean()

In [17]: creditcard_df.isnull().sum()
```

CUST_ID	0
BALANCE	0
BALANCE_FREQUENCY	0
PURCHASES	0
ONEOFF_PURCHASES	0
INSTALLMENTS_PURCHASES	0
CASH_ADVANCE	0
PURCHASES_FREQUENCY	0
ONEOFF_PURCHASES_FREQUENCY	0
PURCHASES_INSTALLMENTS_FREQUENCY	0
CASH_ADVANCE_FREQUENCY	0
CASH_ADVANCE_TRX	0
PURCHASES_TRX	0
CREDIT_LIMIT	1
PAYMENTS	0

MINI CHALLENGE #3:

- Fill out missing elements in the "CREDIT\_LIMIT" column
- Double check and make sure that no missing elements are present

```
In [18]: creditcard_df.loc[(creditcard_df['CREDIT_LIMIT'].isnull() == True),  
                           'CREDIT_LIMIT'] = creditcard_df['CREDIT_LIMIT'].mean()
```

```
In [19]: creditcard_df.isnull().sum()
```

```
CUST_ID      0
BALANCE      0
BALANCE_FREQUENCY  0
PURCHASES    0
ONEOFF_PURCHASES  0
INSTALLMENTS_PURCHASES  0
CASH_ADVANCE  0
PURCHASES_FREQUENCY  0
ONEOFF_PURCHASES_FREQUENCY  0
PURCHASES_INSTALLMENTS_FREQUENCY  0
CASH_ADVANCE_FREQUENCY  0
CASH_ADVANCE_TRX  0
PURCHASES_TRX  0
CREDIT_LIMIT  0
PAYMENTS     0
MINIMUM_PAYMENTS  0
PRC_FULL_PAYMENT  0
TENURE       0
dtype: int64
```

```
In [20]: # Let's see if we have duplicated entries in the data
creditcard_df.duplicated().sum()
```

```
0
```

#### MINI CHALLENGE #4:

- Drop Customer ID column 'CUST\_ID' and make sure that the column has been removed from the dataframe

```
In [21]: creditcard_df.drop('CUST_ID', axis = 1, inplace = True)
```

```
In [22]: creditcard_df
```

	BALANCE	BALANCE_FREQUENCY	PURCHASES	ONEOFF_PURCHASES	INSTALLMENTS_PURCHASES	CASH_ADVANCE
0	40.900749	0.818182	95.40	0.00	95.40	0.000000
1	3202.467416	0.909091	0.00	0.00	0.00	6442.945483
2	2495.148862	1.000000	773.17	773.17	0.00	0.000000
3	1666.670542	0.636364	1499.00	1499.00	0.00	205.788017
4	817.714335	1.000000	16.00	16.00	0.00	0.000000
...	...	...	...	...	...	...
8945	28.493517	1.000000	291.12	0.00	291.12	0.000000
8946	19.183215	1.000000	300.00	0.00	300.00	0.000000
8947	23.398673	0.833333	144.40	0.00	144.40	0.000000

```
8947 23.398673 0.833333 144.40 0.00 144.40 0.000000
8948 13.457564 0.833333 0.00 0.00 0.00 36.558778
8949 372.708075 0.666667 1093.25 1093.25 0.00 127.040008

8950 rows x 7 columns

In [23]: n = len(creditcard_df.columns)
n

17

In [24]: creditcard_df.columns

Index(['BALANCE', 'BALANCE_FREQUENCY', 'PURCHASES', 'ONEOFF_PURCHASES',
      'INSTALLMENTS_PURCHASES', 'CASH_ADVANCE', 'PURCHASES_FREQUENCY',
      'ONEOFF_PURCHASES_FREQUENCY', 'PURCHASES_INSTALLMENTS_FREQUENCY',
      'CASH_ADVANCE_FREQUENCY', 'CASH_ADVANCE_TRX', 'PURCHASES_TRX',
      'CREDIT_LIMIT', 'PAYMENTS', 'MINIMUM_PAYMENTS', 'PRC_FULL_PAYMENT',
      'TENURE'],
      dtype='object')
```


```
In [ ]: # distplot combines the matplotlib.hist function with seaborn kdeplot()
        # KDE Plot represents the Kernel Density Estimate
        # KDE is used for visualizing the Probability Density of a continuous variable.
        # KDE demonstrates the probability density at different values in a continuous variable.




        # Mean of balance is $1500
        # 'Balance_Frequency' for most customers is updated frequently ~1
        # For 'PURCHASES_FREQUENCY', there are two distinct group of customers
        # For 'ONEOFF_PURCHASES_FREQUENCY' and 'PURCHASES_INSTALLMENT_FREQUENCY' most users don't do one off purchase
        # Very small number of customers pay their balance in full 'PRC_FULL_PAYMENT'~0
        # Credit limit average is around $4500
        # Most customers are ~11 years tenure
```



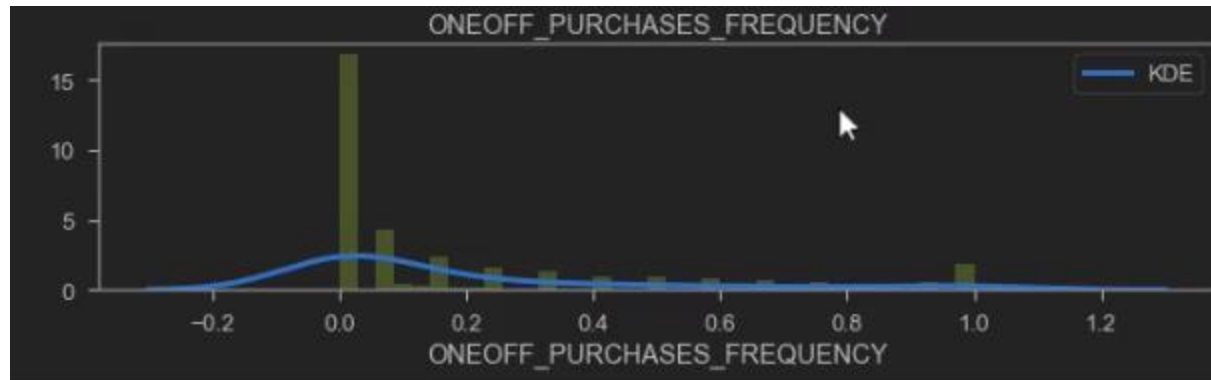


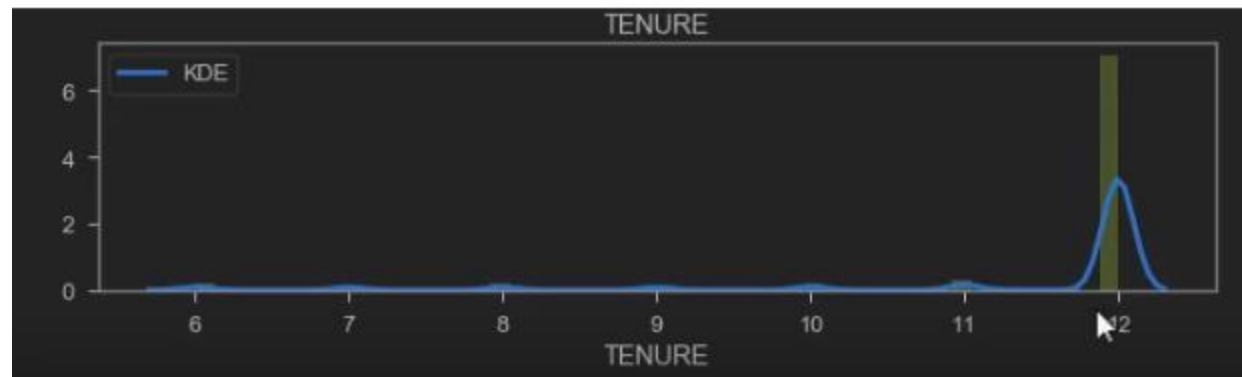
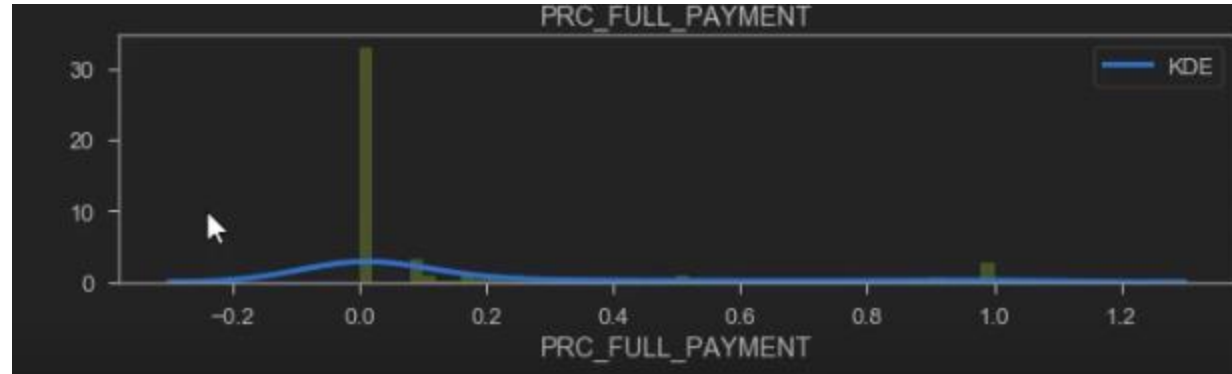
```
plt.figure(figsize=(10,50))
for i in range(len(creditcard_df.columns)):
    plt.subplot(17, 1, i+1)
    sns.distplot(creditcard_df[creditcard_df.columns[i]],
                  kde_kws={"color": "b", "lw": 3, "label": "KDE"}, hist_kws={"color": "g"})
    plt.title(creditcard_df.columns[i])

plt.tight_layout()
```









#### MINI CHALLENGE #5:

- Obtain the correlation matrix between features

```
In [37]: correlations = creditcard_df.corr()  
f, ax = plt.subplots(figsize = (20,10))  
sns.heatmap(correlations, annot = True)
```






## TASK #4: UNDERSTAND THE THEORY AND INTUITION BEHIND K-MEANS


### K-MEANS INTUITION

- K-means is an unsupervised learning algorithm (clustering).
- K-means works by grouping some data points together (clustering) in an unsupervised fashion.
- The algorithm groups observations with similar attribute values together by measuring the Euclidian distance between points.



## K-MEANS ALGORITHM STEPS

1. Choose number of clusters "K"
2. Select random K points that are going to be the centroids for each cluster
3. Assign each data point to the nearest centroid, doing so will enable us to create "K" number of clusters
4. Calculate a new centroid for each cluster
5. Reassign each data point to the new closest centroid
6. Go to step 4 and repeat.



MINI CHALLENGE #6:


- Which of the following conditions could terminate the K-means clustering algorithm? (choose 2)
  - K-means terminates after a fixed number of iterations is reached (correct)
  - K-means terminates when the number of clusters does not increase between iterations (wrong)
  - K-means terminates when the centroid locations do not change between iterations (correct)

## TASK #5: LEARN HOW TO OBTAIN THE OPTIMAL NUMBER OF CLUSTERS (ELBOW METHOD)

### HOW TO SELECT THE OPTIMAL NUMBER OF CLUSTERS (K)? “ELBOW METHOD”

*Within Cluster Sum of Squares (WCSS)*

$$= \sum_{P_i \text{ in Cluster 1}} \text{distance}(P_i, C_1)^2 + \sum_{P_i \text{ in Cluster 2}} \text{distance}(P_i, C_2)^2 + \sum_{P_i \text{ in Cluster 3}} \text{distance}(P_i, C_3)^2$$



# HOW TO SELECT THE OPTIMAL NUMBER OF CLUSTERS (K)?

## “ELBOW METHOD”

$$\text{Within Cluster Sum of Squares (WCSS)} = \sum_{P_i \text{ in Cluster } 1} \text{distance}(P_i, C_1)^2$$

**NUMBER OF CLUSTERS (K) = 1**






# HOW TO SELECT THE OPTIMAL NUMBER OF CLUSTERS (K)?

## “ELBOW METHOD”

$$\text{Within Cluster Sum of Squares (WCSS)} = \sum_{P_i \text{ in Cluster 1}} \text{distance}(P_i, C_1)^2 + \sum_{P_i \text{ in Cluster 2}} \text{distance}(P_i, C_2)^2$$

**NUMBER OF CLUSTERS (K) = 2**




# HOW TO SELECT THE OPTIMAL NUMBER OF CLUSTERS (K)?

## “ELBOW METHOD”


Within Cluster Sum of Squares (WCSS)

$$= \sum_{P_i \text{ in Cluster 1}} \text{distance}(P_i, C_1)^2 + \sum_{P_i \text{ in Cluster 2}} \text{distance}(P_i, C_2)^2 + \sum_{P_i \text{ in Cluster 3}} \text{distance}(P_i, C_3)^2$$

**NUMBER OF CLUSTERS (K) = 3**



## HOW TO SELECT THE OPTIMAL NUMBER OF CLUSTERS (K)? “ELBOW METHOD”



Source: [https://commons.wikimedia.org/wiki/File:Tennis\\_Elbow\\_Illustration.jpg](https://commons.wikimedia.org/wiki/File:Tennis_Elbow_Illustration.jpg)

## TASK #6: FIND THE OPTIMAL NUMBER OF CLUSTERS USING ELBOW METHOD

- The elbow method is a heuristic method of interpretation and validation of consistency within cluster analysis designed to help find the appropriate number of clusters in a dataset.
- If the line chart looks like an arm, then the "elbow" on the arm is the value of k that is the best.
- Source:
  - [https://en.wikipedia.org/wiki/Elbow\\_method\\_\(clustering\)](https://en.wikipedia.org/wiki/Elbow_method_(clustering))
  - <https://www.geeksforgeeks.org/elbow-method-for-optimal-value-of-k-in-kmeans/>

```
In [38]: # Let's scale the data first
        scaler = StandardScaler()
        creditcard_df_scaled = scaler.fit_transform(creditcard_df)
```

```
In [39]: creditcard_df_scaled.shape
```

```
(8950, 17)
```

```
In [40]: creditcard_df_scaled
```

```
array([[ -0.73198937, -0.24943448, -0.42489974, ..., -0.31096755,
        -0.52555097,  0.36067954],
       [ 0.78696085,  0.13432467, -0.46955188, ...,  0.08931021,
        0.2342269 ,  0.36067954],
       [ 0.44713513,  0.51808382, -0.10766823, ..., -0.10166318,
        -0.52555097,  0.36067954],
       ...,
       [-0.7403981 , -0.18547673, -0.40196519, ..., -0.33546549,
        0.32919999, -4.12276757],
       [-0.74517423, -0.18547673, -0.46955188, ..., -0.34690648,
        0.32919999, -4.12276757],
       [-0.57257511, -0.88903307,  0.04214581, ..., -0.33294642,
        -0.52555097, -4.12276757]])
```

```
In [ ]: # Index(['BALANCE', 'BALANCE_FREQUENCY', 'PURCHASES', 'ONEOFF_PURCHASES',
#             'INSTALLMENTS_PURCHASES', 'CASH_ADVANCE', 'PURCHASES_FREQUENCY',
#             'ONEOFF_PURCHASES_FREQUENCY', 'PURCHASES_INSTALLMENTS_FREQUENCY',
#             'CASH_ADVANCE_FREQUENCY', 'CASH_ADVANCE_TRX', 'PURCHASES_TRX',
#             'CREDIT_LIMIT', 'PAYMENTS', 'MINIMUM_PAYMENTS', 'PRC_FULL_PAYMENT',
#             'TENURE'], dtype='object')
```

```
# From this we can observe that, 4th cluster seems to be forming the elbow of the curve.
# However, the values does not reduce linearly until 8th cluster.
# Let's choose the number of clusters to be 7 or 8.
```

```


In [42]: # Index(['BALANCE', 'BALANCE_FREQUENCY', 'PURCHASES', 'ONEOFF_PURCHASES',
#             'INSTALLMENTS_PURCHASES', 'CASH_ADVANCE', 'PURCHASES_FREQUENCY',
#             'ONEOFF_PURCHASES_FREQUENCY', 'PURCHASES_INSTALLMENTS_FREQUENCY',
#             'CASH_ADVANCE_FREQUENCY', 'CASH_ADVANCE_TRX', 'PURCHASES_TRX',
#             'CREDIT_LIMIT', 'PAYMENTS', 'MINIMUM_PAYMENTS', 'PRC_FULL_PAYMENT',
#             'TENURE'], dtype='object')

scores_1 = []
range_values = range(1,20)
for i in range_values:
    kmeans = KMeans(n_clusters = i)
    kmeans.fit(creditcard_df_scaled)
    scores_1.append(kmeans.inertia_)

plt.plot(scores_1, 'bx-')
# From this we can observe that, 4th cluster seems to be forming the elbow of the curve.
# However, the values does not reduce linearly until 8th cluster.
# Let's choose the number of clusters to be 7 or 8.

[<matplotlib.lines.Line2D at 0x2628fb7cac8>]

```





MINI CHALLENGE #7:

- Let's assume that our data only consists of the first 7 columns of "creditcard\_df\_scaled", what is the optimal number of clusters would be in this case? modify the code and rerun the cells.

```
In [43]: creditcard_df_scaled[:, :7].shape
```

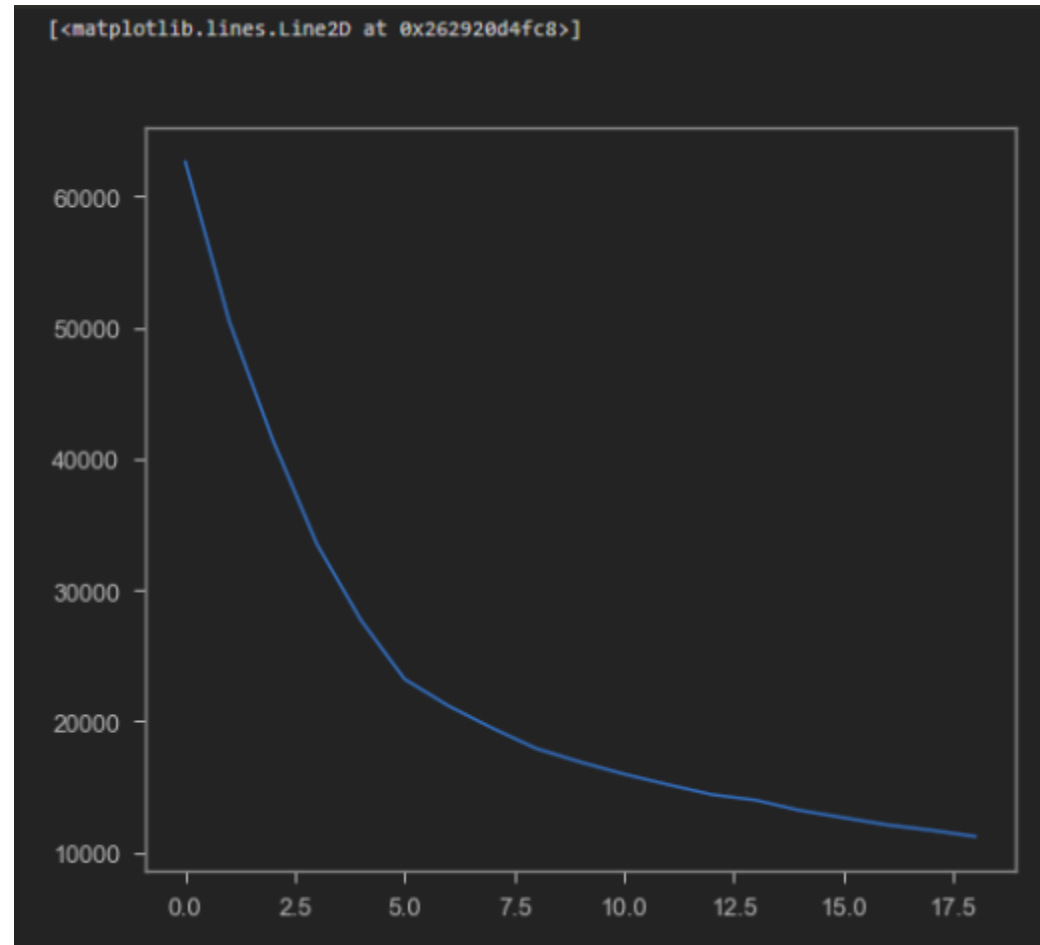
```
(8950, 7)
```

```
In [44]: scores_1 = []
range_values = range(1,20)
for i in range_values:
    kmeans = KMeans(n_clusters = i)
    kmeans.fit(creditcard_df_scaled[:, :7])
    scores_1.append(kmeans.inertia_)

plt.plot(scores_1, 'bx-')
```

```
[<matplotlib.lines.Line2D at 0x262920d4fc8>]
```





## TASK #7: APPLY K-MEANS METHOD

```
In [45]: kmeans = KMeans(7)
kmeans.fit(creditcard_df_scaled)
labels = kmeans.labels_ # Labels (cluster) associated to each data point
```

```
In [46]: kmeans.cluster_centers_.shape
```

```
(7, 17)
```

```
In [47]: cluster_centers = pd.DataFrame(data = kmeans.cluster_centers_, columns = [creditcard_df.columns])
cluster_centers
```

	BALANCE	BALANCE_FREQUENCY	PURCHASES	ONEOFF_PURCHASES	INSTALLMENTS_PURCHASES	CASH_ADVANCE
0	1.666473	0.392099	-0.205327	-0.149913	-0.210162	1.990753
1	-0.701872	-2.134325	-0.306924	-0.230292	-0.302515	-0.323078
2	-0.367555	0.330562	-0.039840	-0.234950	0.337266	-0.368099
3	-0.335506	-0.348076	-0.284525	-0.208973	-0.288475	0.065539
4	0.007813	0.402983	-0.343915	-0.225214	-0.399316	-0.104212
5	0.126801	0.429730	0.939029	0.895888	0.574411	-0.309125
6	1.430238	0.419467	6.915048	6.083034	5.172266	0.038778

```
In [48]: # In order to understand what these numbers mean, let's perform inverse transformation
cluster_centers = scaler.inverse_transform(cluster_centers)
cluster_centers = pd.DataFrame(data = cluster_centers, columns = [creditcard_df.columns])
cluster_centers

# First Customers cluster (Transactors): Those are customers who pay least amount of intrerest charges and
# Second customers cluster (revolvers) who use credit card as a loan (most lucrative sector): highest balance
# Third customer cluster (VIP/Prime): high credit limit $16K and highest percentage of full payment, target
# Fourth customer cluster (low tenure): these are customers with low tenure (7 years), low balance
```

	BALANCE	BALANCE_FREQUENCY	PURCHASES	ONEOFF_PURCHASES	INSTALLMENTS_PURCHASES	CASH_ADVANCE
0	5033.096672	0.970155	564.520447	343.613277	221.020895	5153.573564
1	103.587241	0.371669	347.456361	210.199629	137.506773	301.361116
2	799.439665	0.955578	918.085545	202.468866	716.053491	206.951098
3	866.148306	0.794815	395.311749	245.585564	150.203132	1116.308792
4	1580.736068	0.972734	268.425678	218.628807	49.971063	760.334088
5	1828.399674	0.979070	3009.455142	2079.427650	930.500678	330.620880
6	4541.393882	0.976638	15777.311395	10689.027791	5088.283605	1060.190695

```
In [49]: labels.shape # Labels associated to each data point
```

```
(8950,)
```

```
In [50]: labels.max()
```

```
6
```

```
In [51]: labels.min()
```

```
0
```

```
In [52]: y_kmeans = kmeans.fit_predict(creditcard_df_scaled)
y_kmeans
```

```
array([2, 6, 5, ..., 3, 3, 3])
```

```
In [53]: # concatenate the clusters labels to our original dataframe
creditcard_df_cluster = pd.concat([creditcard_df, pd.DataFrame({'cluster': labels})], axis = 1)
creditcard_df_cluster.head()
```

	BALANCE	BALANCE_FREQUENCY	PURCHASES	ONEOFF_PURCHASES	INSTALLMENTS_PURCHASES	CASH_ADVANCE
0	40.900749	0.818182	95.40	0.00	95.4	0.000000
1	3202.467416	0.909091	0.00	0.00	0.0	6442.945483
2	2495.148862	1.000000	773.17	773.17	0.0	0.000000
3	1666.670542	0.636364	1499.00	1499.00	0.0	205.788017
4	817.714335	1.000000	16.00	16.00	0.0	0.000000

PURCHASES_FREQUENCY	ONEOFF_PURCHASES_FREQUENCY	PURCHASES_INSTALLMENTS_FREQUENCY	C/
0.307019	0.145697	0.203860	0.5
0.270757	0.074792	0.188993	0.0
0.883916	0.095855	0.830697	0.0
0.410589	0.121144	0.272729	0.1
0.165154	0.102107	0.065464	0.1
0.928932	0.761065	0.576983	0.0
0.928101	0.763090	0.781501	0.0

CASH_ADVANCE_FREQUENCY	CASH_ADVANCE_TRX	PURCHASES_TRX	CREDIT_LIMIT	PAYMENTS	MINIMUM_P
0.517064	16.365772	8.709172	8160.463697	4149.870001	2152.885952
0.030637	0.677338	4.355518	3865.955724	1149.108580	263.988609
0.039554	0.775550	19.032274	3482.112568	1091.026929	827.537625
0.196000	3.233704	5.125596	2468.226470	602.104087	376.247870
0.152124	2.964097	3.161211	3399.161094	1012.580763	827.406808
0.053020	1.059937	44.458991	7047.418985	2847.821449	730.293998
0.085271	2.988372	130.197674	12493.023256	15581.496801	3383.304083

MINIMUM_PAYMENTS	PRC_FULL_PAYMENT	TENURE
2152.885952	0.039307	11.610738
263.988609	0.240000	11.786015
827.537625	0.243576	11.854768
376.247870	0.157487	7.243243
827.406808	0.021266	11.881732
730.293998	0.287217	11.929022
3383.304083	0.394721	11.965116

```
In [53]: # concatenate the clusters labels to our original dataframe
creditcard_df_cluster = pd.concat([creditcard_df, pd.DataFrame({'cluster': labels})], axis = 1)
creditcard_df_cluster.head()
```


I_ADVANCE_TRX	PURCHASES_TRX	CREDIT_LIMIT	PAYMENTS	MINIMUM_PAYMENTS	PRC_FULL_PAYMENT	TENURE	cluster
2		1000.0	201.802084	139.509787	0.000000	12	4
0		7000.0	4103.032597	1072.340217	0.222222	12	0
12		7500.0	622.066742	627.284787	0.000000	12	5
1		7500.0	0.000000	864.206542	0.000000	12	4
1		1200.0	678.334763	244.791237	0.000000	12	4

```

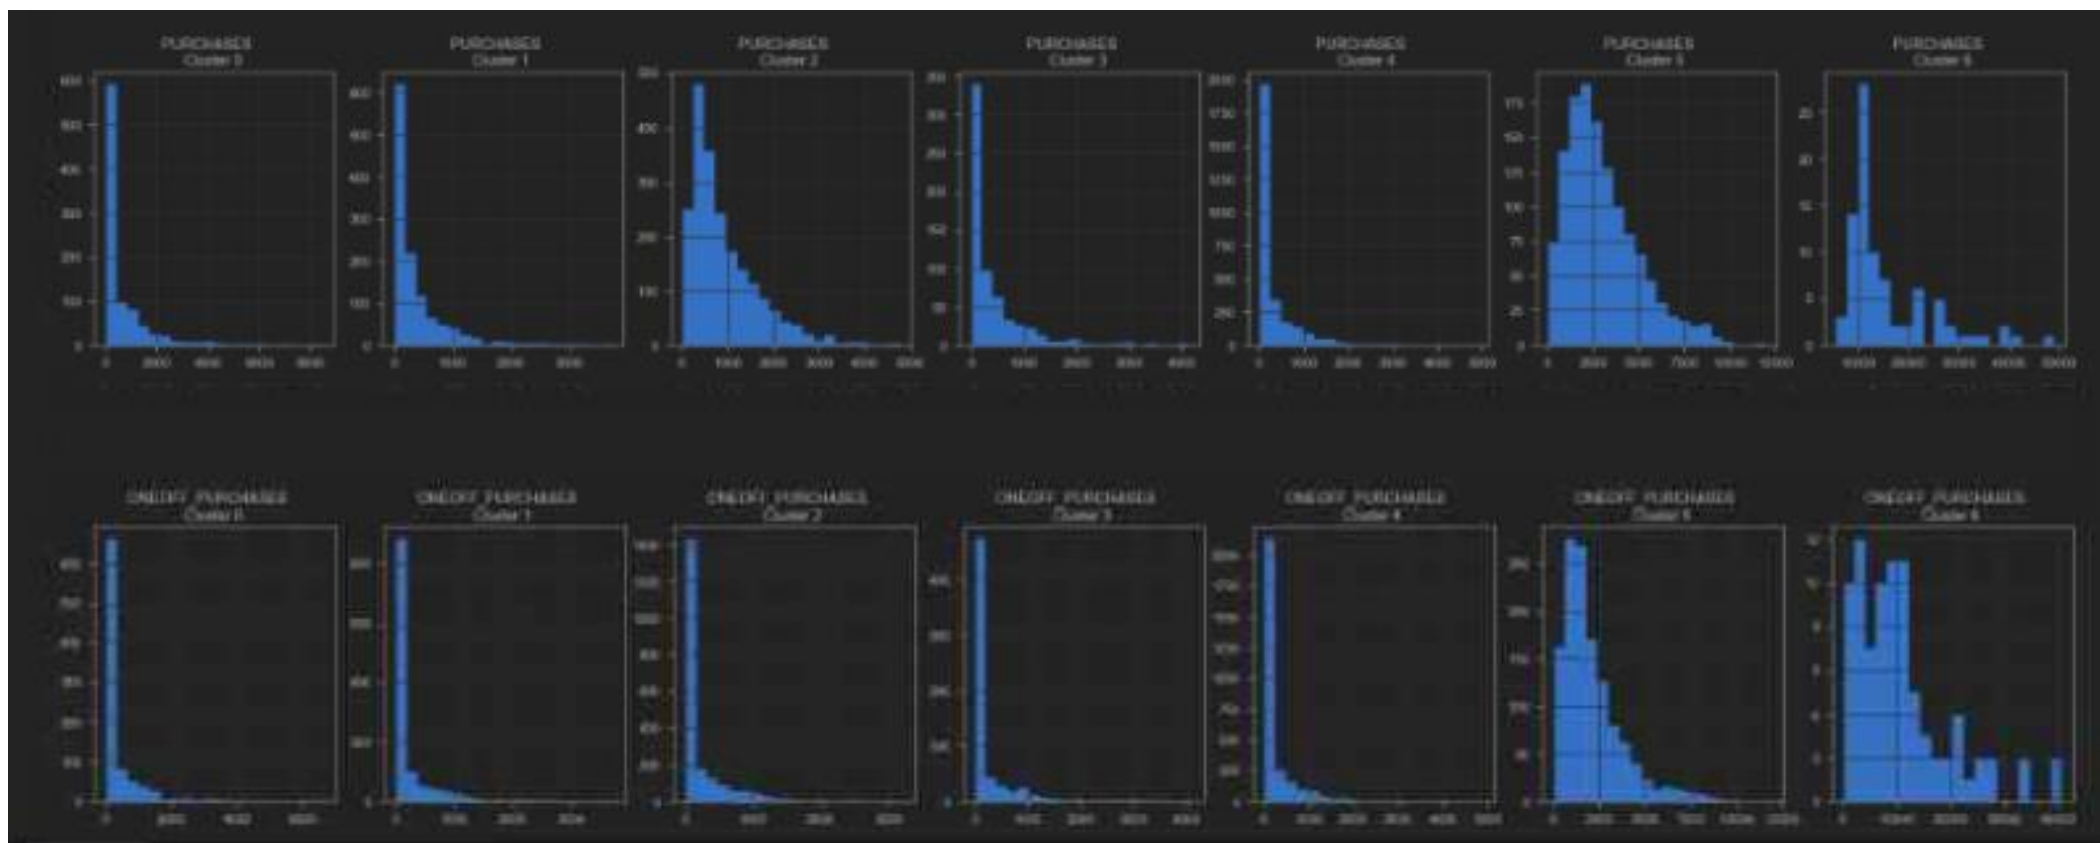
In [54]: # Plot the histogram of various clusters
for i in creditcard_df.columns:
    plt.figure(figsize = (35, 5))
    for j in range(7):
        plt.subplot(1,7,j+1)
        cluster = creditcard_df_cluster[creditcard_df_cluster['cluster'] == j]
        cluster[i].hist(bins = 20)
        plt.title('{} \nCluster {}'.format(i,j))

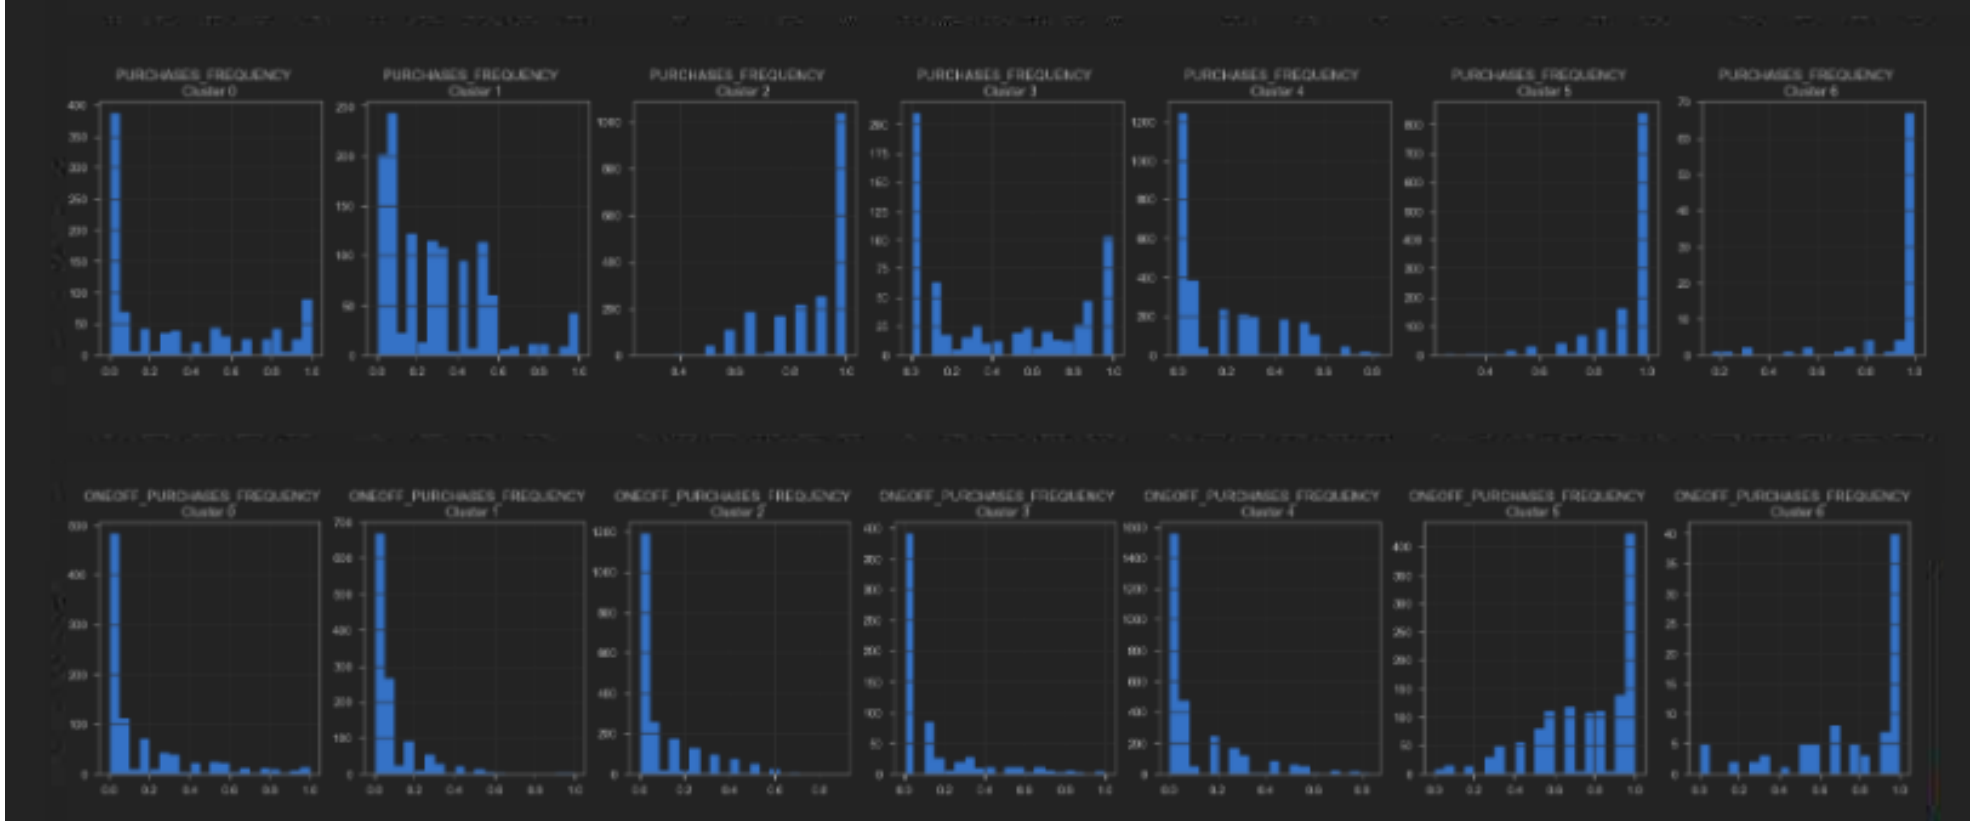
plt.show()

```









### MINI CHALLENGE #8:

- Repeat the same procedure with 8 or 5 or 4 clusters instead of 7

```
In [56]: kmeans = KMeans(5)
         kmeans.fit(creditcard_df_scaled)
         labels = kmeans.labels_ # Labels (cluster) associated to each data point
```

```
In [57]: kmeans.cluster_centers_.shape
```

(5, 17)

```
In [58]: cluster_centers = pd.DataFrame(data = kmeans.cluster_centers_, columns = [creditcard_df.columns])
         cluster_centers
```

	BALANCE	BALANCE_FREQUENCY	PURCHASES	ONEOFF_PURCHASES	INSTALLMENTS_PURCHASES	CASH_ADVANCE
0	-0.242188	-0.335038	-0.339548	-0.221313	-0.396128	-0.168494
1	0.157739	0.429926	0.895701	0.835877	0.582143	-0.298657
2	1.430238	0.419467	6.915048	6.083034	5.172266	0.038778
3	-0.443442	0.103191	-0.099301	-0.257555	0.238336	-0.381977
4	1.481094	0.382972	-0.235548	-0.172446	-0.240230	1.760672

```
cluster_centers = scaler.inverse_transform(cluster_centers)
cluster_centers = pd.DataFrame(data = cluster_centers, columns = [creditcard_df.columns])
cluster_centers
```

*# First Customers cluster (Transactors): Those are customers who pay least amount of intrerest charges and i*  
*# Second customers cluster (revolvers) who use credit card as a loan (most lucrative sector): highest balan*  
*# Third customer cluster (VIP/Prime): high credit limit \$16K and highest percentage of full payment, target*  
*# Fourth customer cluster (Low tenure): these are customers with low tenure (7 years), Low balance*

	CASH_ADVANCE_TRX	PURCHASES_TRX	CREDIT_LIMIT	PAYMENTS	MINIMUM_PAYMENTS	PRC_FULL_PAYMENT	TENURE
2.243003	2.941476	3350.334466	1005.991021	602.097184	0.069941	11.456234	
1.153127	43.511143	6960.026796	2776.202091	834.876950	0.264311	11.888569	
2.988372	130.197674	12493.023256	15581.496801	3383.304083	0.394721	11.965116	
0.708299	16.563071	3198.301375	929.237052	674.645925	0.273139	11.453942	
14.734334	7.462489	7585.882211	3633.650163	2021.387424	0.036197	11.374228	

```
In [60]: labels.shape # Labels associated to each data point
```

```
(8950,)
```

```
In [61]: labels.max()
```

```
4
```

```
In [62]: labels.min()
```

```
0
```

```
In [63]: y_kmeans = kmeans.fit_predict(creditcard_df_scaled)
y_kmeans
```

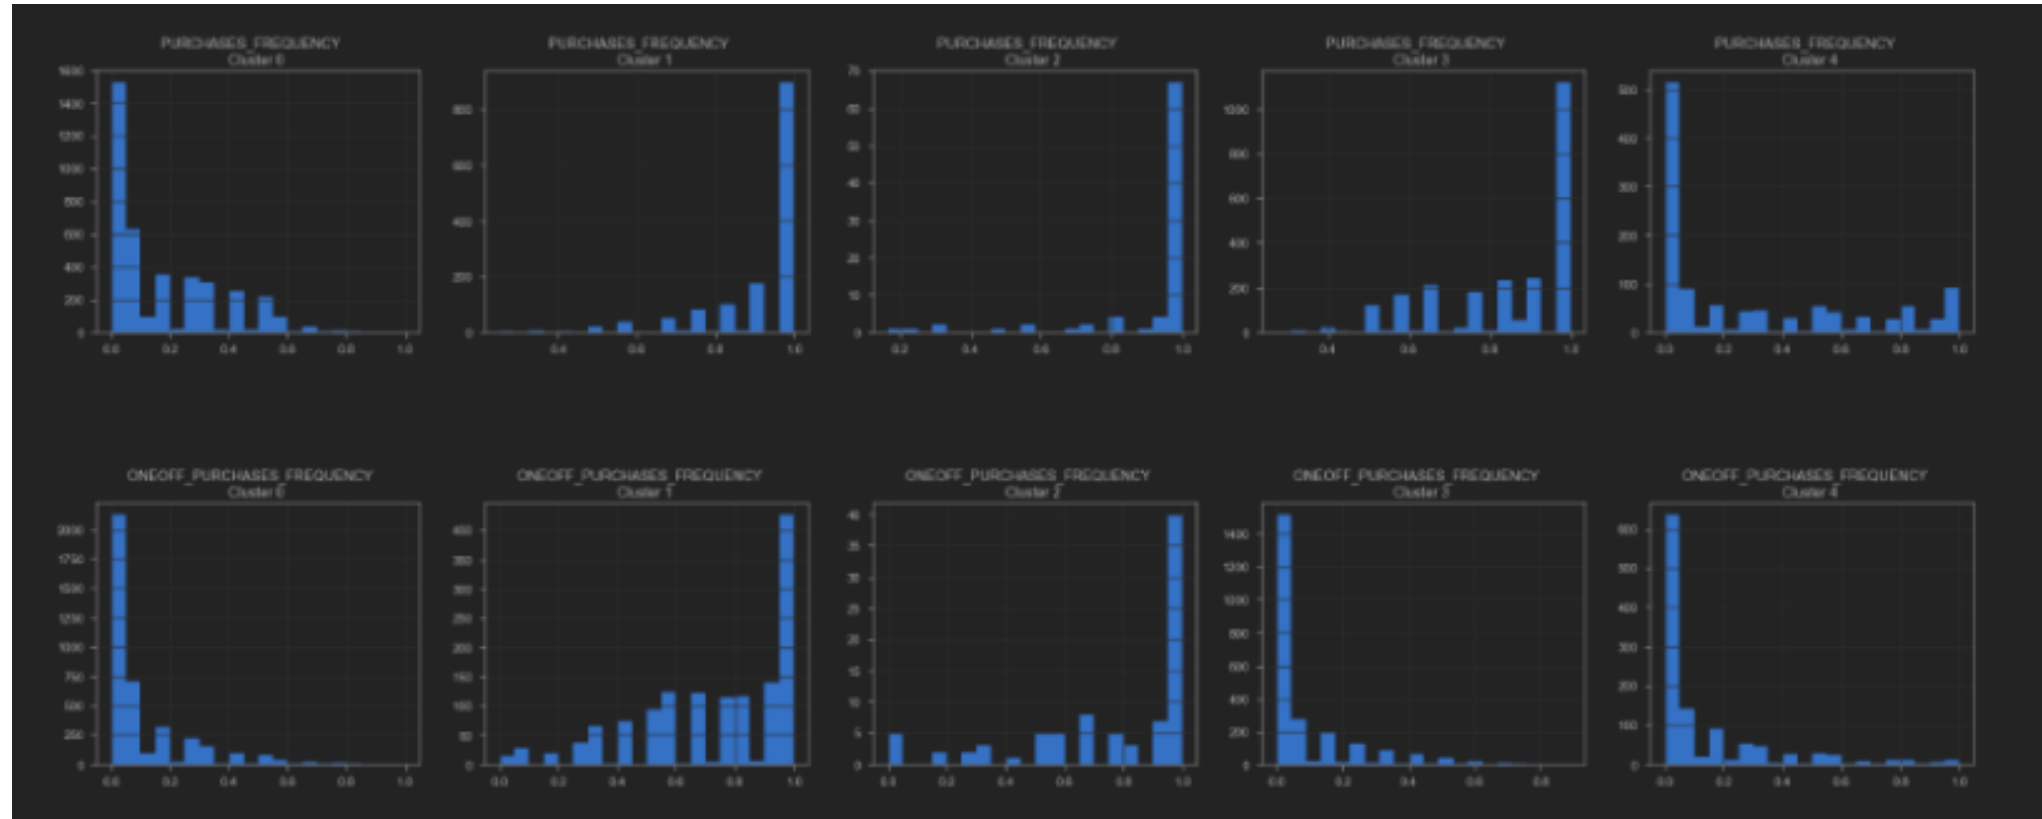
```
array([4, 2, 1, ..., 3, 4, 4])
```

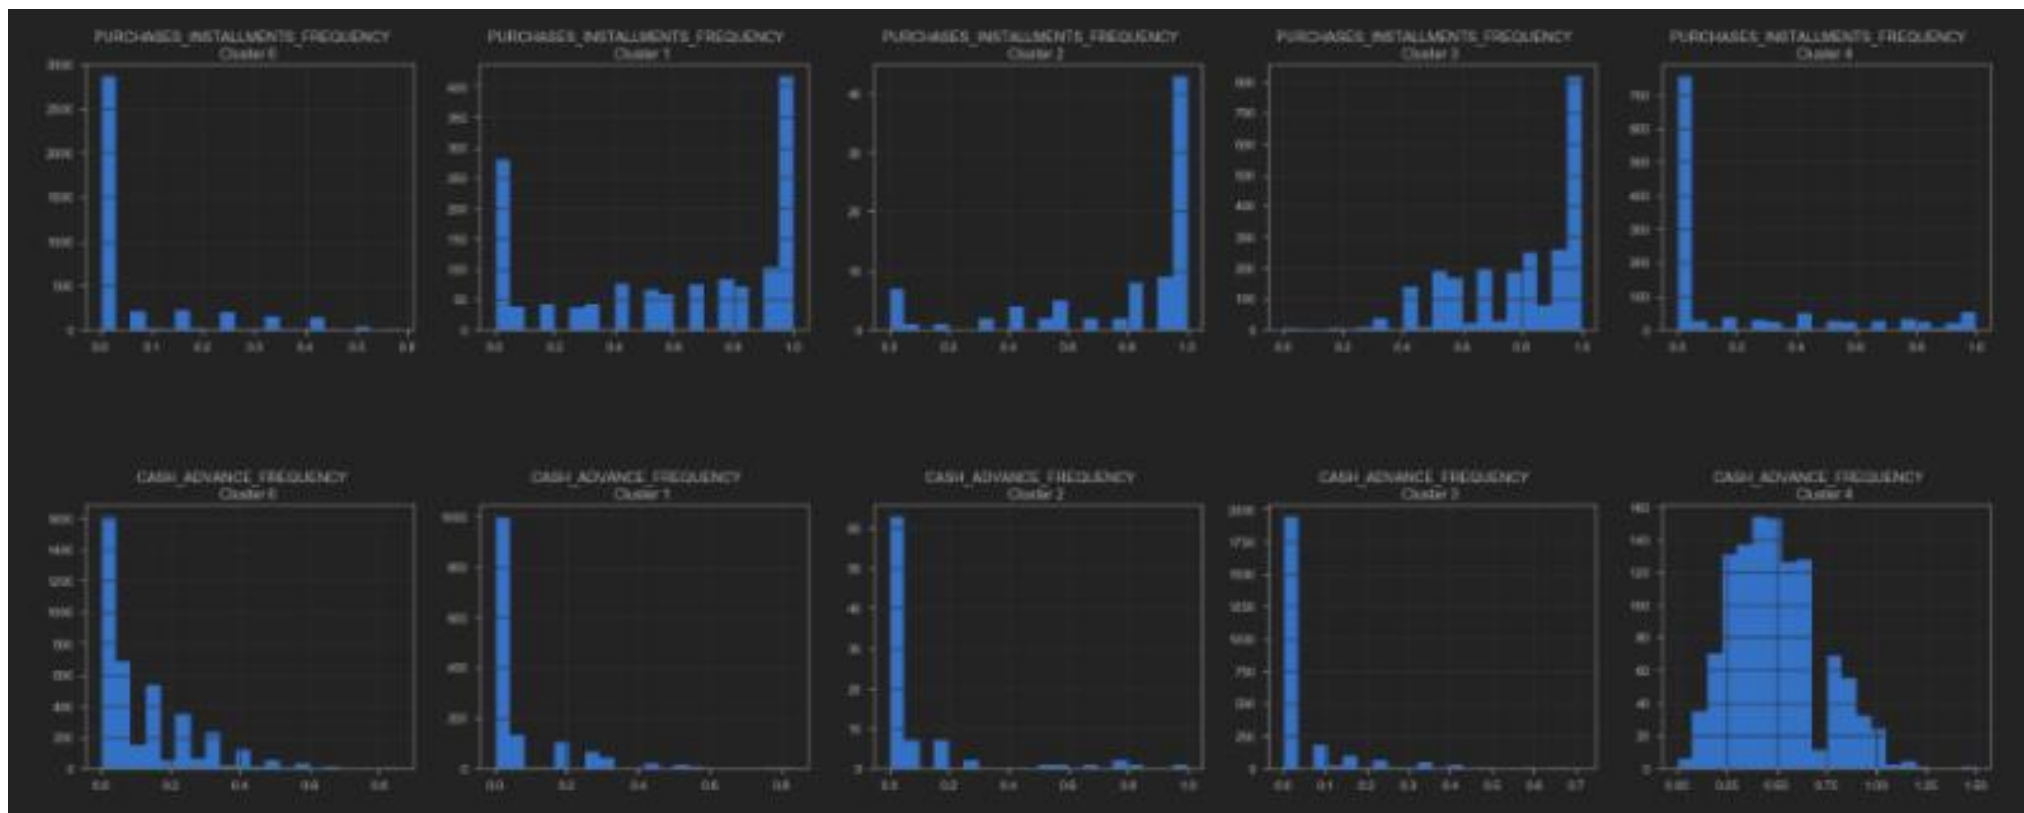
```
In [64]: # concatenate the clusters labels to our original dataframe
creditcard_df_cluster = pd.concat([creditcard_df, pd.DataFrame({'cluster': labels})], axis = 1)
creditcard_df_cluster.head()
```

I_ADVANCE_TRX	PURCHASES_TRX	CREDIT_LIMIT	PAYMENTS	MINIMUM_PAYMENTS	PRC_FULL_PAYMENT	TENURE	cluster
2		1000.0	201.802084	139.509787	0.000000	12	0
0		7000.0	4103.032597	1072.340217	0.222222	12	4
12		7500.0	622.066742	627.284787	0.000000	12	1
1		7500.0	0.000000	864.206542	0.000000	12	0
1		1200.0	678.334763	244.791237	0.000000	12	0

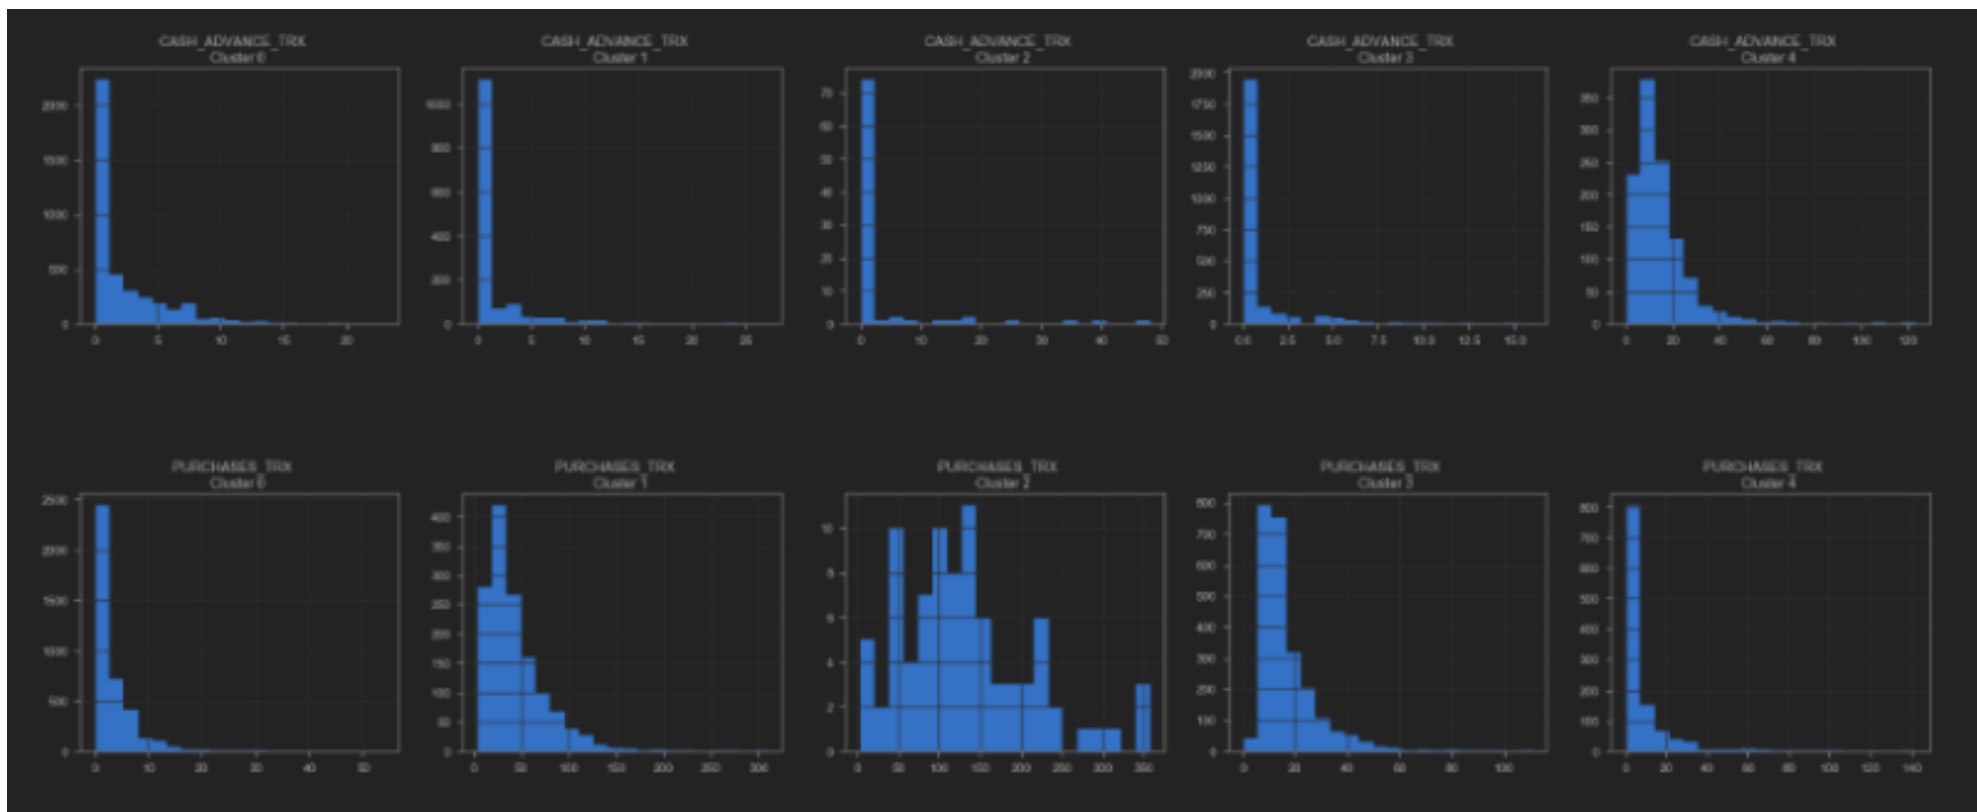
```
In [*]: # Plot the histogram of various clusters
for i in creditcard_df.columns:
    plt.figure(figsize = (35, 5))
    for j in range(5):
        plt.subplot(1,5,j+1)
        cluster = creditcard_df_cluster[creditcard_df_cluster['cluster'] == j]
        cluster[i].hist(bins = 20)
        plt.title('{} \nCluster {}'.format(i,j))

plt.show()
```









# TASK 8: APPLY PRINCIPAL COMPONENT ANALYSIS AND VISUALIZE THE RESULTS

## PRINCIPAL COMPONENT ANALYSIS (PCA)

- PCA is an unsupervised machine learning algorithm.
- PCA performs dimensionality reductions while attempting at keeping the original information unchanged.
- PCA works by trying to find a new set of features called components.
- Components are composites of the uncorrelated given input features.

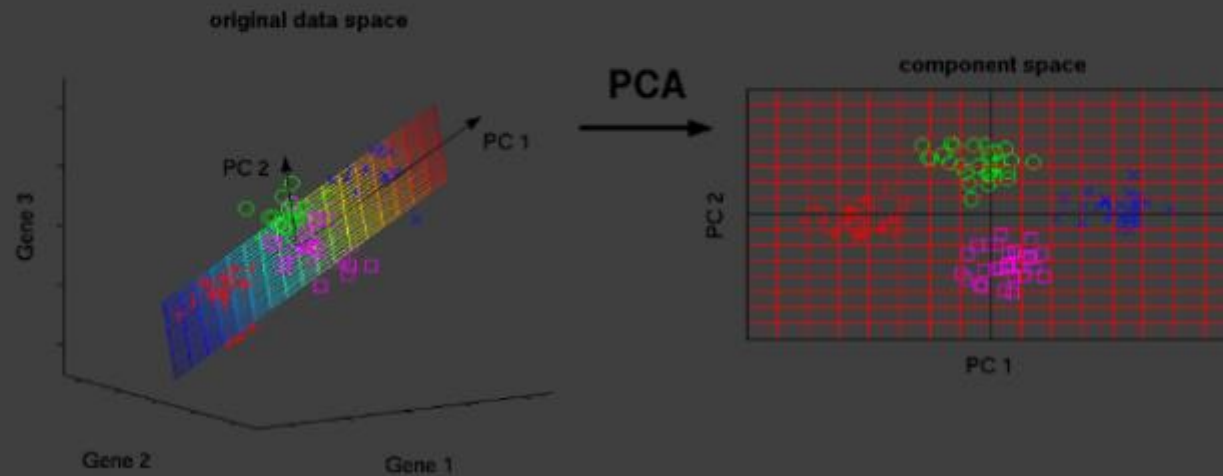


Photo Credit: <http://phdthesis-bioinformatics-maxplanckinstitute-molecularplantphys.matthias-scholz.de/>

```
In [66]: # Obtain the principal components
pca = PCA(n_components=2)
principal_comp = pca.fit_transform(creditcard_df_scaled)
principal_comp
```

```
array([[ -1.68221882, -1.07645695],
       [ -1.1383011 ,  2.5065022 ],
       [  0.96968621, -0.38352992],
       ...,
       [ -0.92619882, -1.81080592],
       [ -2.33654972, -0.65797375],
       [ -0.55642724, -0.40044493]])
```

```
In [67]: # Create a dataframe with the two components
pca_df = pd.DataFrame(data = principal_comp, columns = ['pca1', 'pca2'])
pca_df.head()
```

	pca1	pca2
0	-1.682219	-1.076457
1	-1.138301	2.506502
2	0.969686	-0.383530
3	-0.873628	0.043165
4	-1.599431	-0.688591

```
In [68]: # Concatenate the clusters labels to the dataframe
pca_df = pd.concat([pca_df, pd.DataFrame({'cluster': labels})], axis = 1)
pca_df.head()
```

	pca1	pca2	cluster
0	-1.682219	-1.076457	0
1	-1.138301	2.506502	4
2	0.969686	-0.383530	1
3	-0.873628	0.043165	0
4	-1.599431	-0.688591	0

```
In [69]: plt.figure(figsize=(10,10))
ax = sns.scatterplot(x="pca1", y="pca2", hue = "cluster",
                    data = pca_df, palette = ['red', 'green', 'blue', 'pink', 'yellow'])
plt.show()
```

